Figure 2. Performance attribution for the HG002 Revio chr20 end-to-end benchmark (8 threads)

Lower is better. The language change alone contributed little; profiling and pipeline restructuring produced the gain.

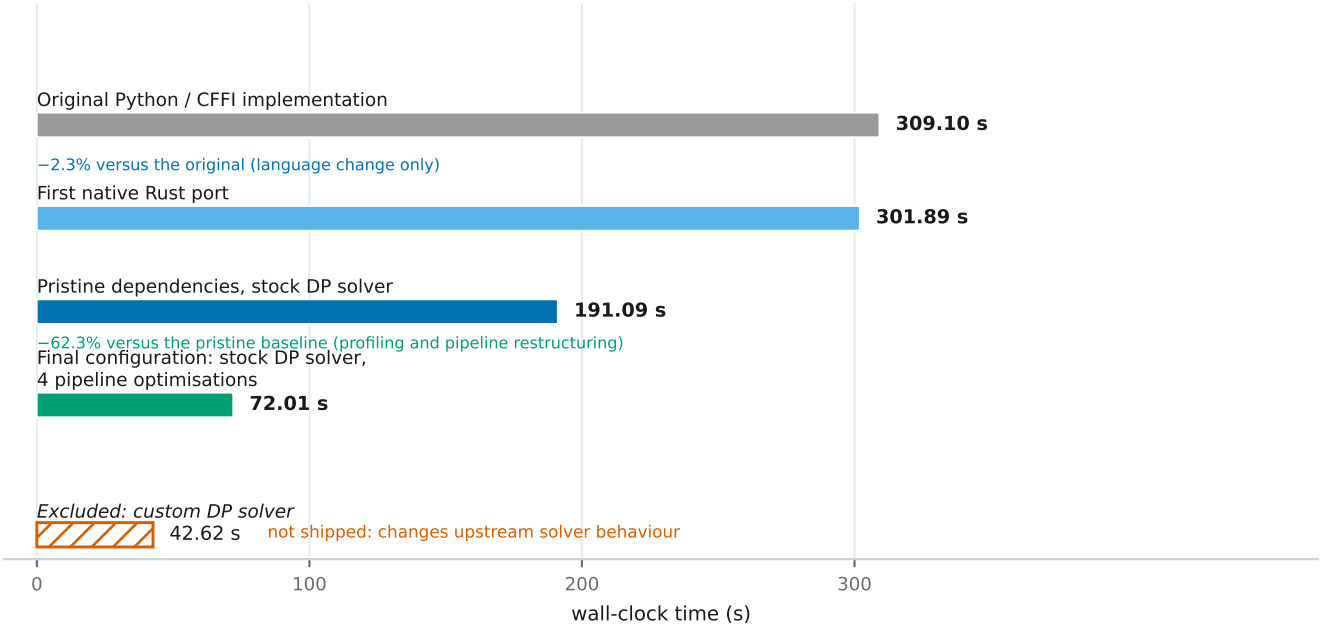

One measurement campaign on the platform of Table 2; peak RSS 1.97 GB (original) → 3.93 GB (final).
Emulated and CUDA timings are functional measurements, not native benchmarks.